

mmn_12_computer_vision

December 13, 2018

1 MMN 12 Computer Vision

1.1 Preparatory Setup

1.1.1 Library Imports

```
In [184]: import numpy as np
import cv2
import matplotlib.pyplot as plt
from mpl_toolkits.mplot3d import Axes3D
import seaborn as sns
import pandas as pd
from sklearn.utils import shuffle
from sklearn.model_selection import train_test_split
from sklearn.neighbors import KNeighborsClassifier
from sklearn.preprocessing import StandardScaler, label_binarize
from sklearn.metrics import roc_curve, auc, classification_report, confusion_matrix
from sklearn.model_selection import GridSearchCV, RandomizedSearchCV
from sklearn.svm import SVC, LinearSVC
from scipy.stats import uniform, expon, randint
from sklearn.decomposition import PCA, TruncatedSVD
from sklearn.manifold import TSNE
import os
%matplotlib inline
```

1.1.2 Plot Style Setup

```
In [2]: # change rcParams so figure aesthetics match jupyterthemes style
# NOTE: or use plt.style.use(['dark_background']) to get dark plots without installing
from jupyterthemes import jtplot
jtplot.style()
```

1.2 Iris Dataset

[link to Iris dataset source website](#)

1.2.1 Overview

The dataset consists of 150 samples. Each sample is assigned to 1 of 3 possible classes, and holds a value for each of the 4 attributes.

Note: this is not the original dataset in neither shape, form, or format. We've done some pre-processing (outside notebook scope) for convenience.

Classes

- Iris Setosa
- Iris Versicolour
- Iris Virginica

Attributes

1. sepal length in cm
2. sepal width in cm
3. petal length in cm
4. petal width in cm

1.2.2 Dataset Exploration

```
In [10]: # load dataset
iris_data = pd.read_csv('./iris_dataset.csv', index_col='id')
```

```
In [11]: # number of samples
iris_data.species.count()
```

```
Out[11]: 150
```

```
In [12]: # showcase a few random samples
iris_data.sample(5)
```

```
Out[12]:
```


	sepal_length	sepal_width	petal_length	petal_width	species
id					
104	6.5	3.0	5.8	2.2	virginica
74	6.4	2.9	4.3	1.3	versicolor
69	5.6	2.5	3.9	1.1	versicolor
30	4.8	3.1	1.6	0.2	setosa
36	5.5	3.5	1.3	0.2	setosa

```
In [6]: # unique classes
pd.unique(iris_data.species)
```

```
Out[6]: array(['setosa', 'versicolor', 'virginica'], dtype=object)
```

```
In [8]: # number of samples in each class
iris_data.species.value_counts().plot.bar()
```

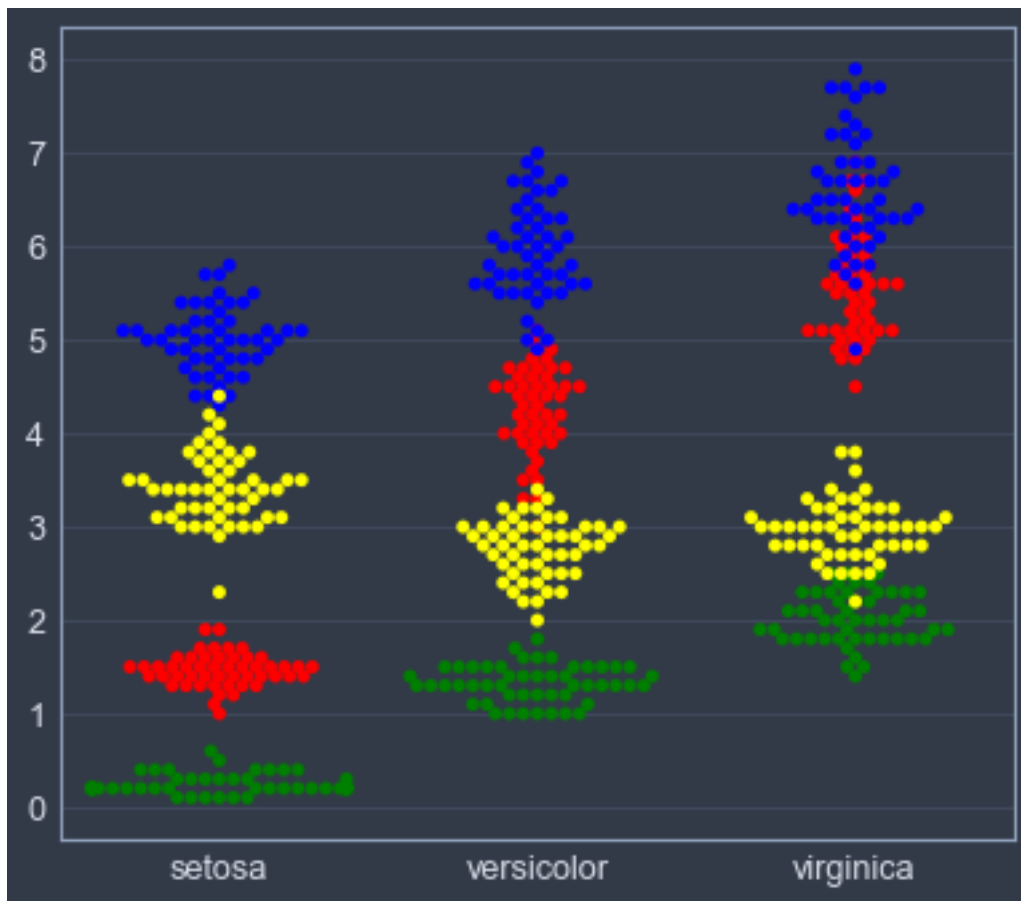
```
Out[8]: <matplotlib.axes._subplots.AxesSubplot at 0x115a50978>
```



```
In [21]: # x_name = 'sepal_width'; y_name = 'sepal_length'
# for species in pd.unique(iris_data.species):
#     x = iris_data[iris_data.species==species][x_name]
#     y = iris_data[iris_data.species==species][y_name]
#     plt.scatter(x,y, label=species)
# plt.xlabel(x_name); plt.ylabel(y_name); plt.legend()

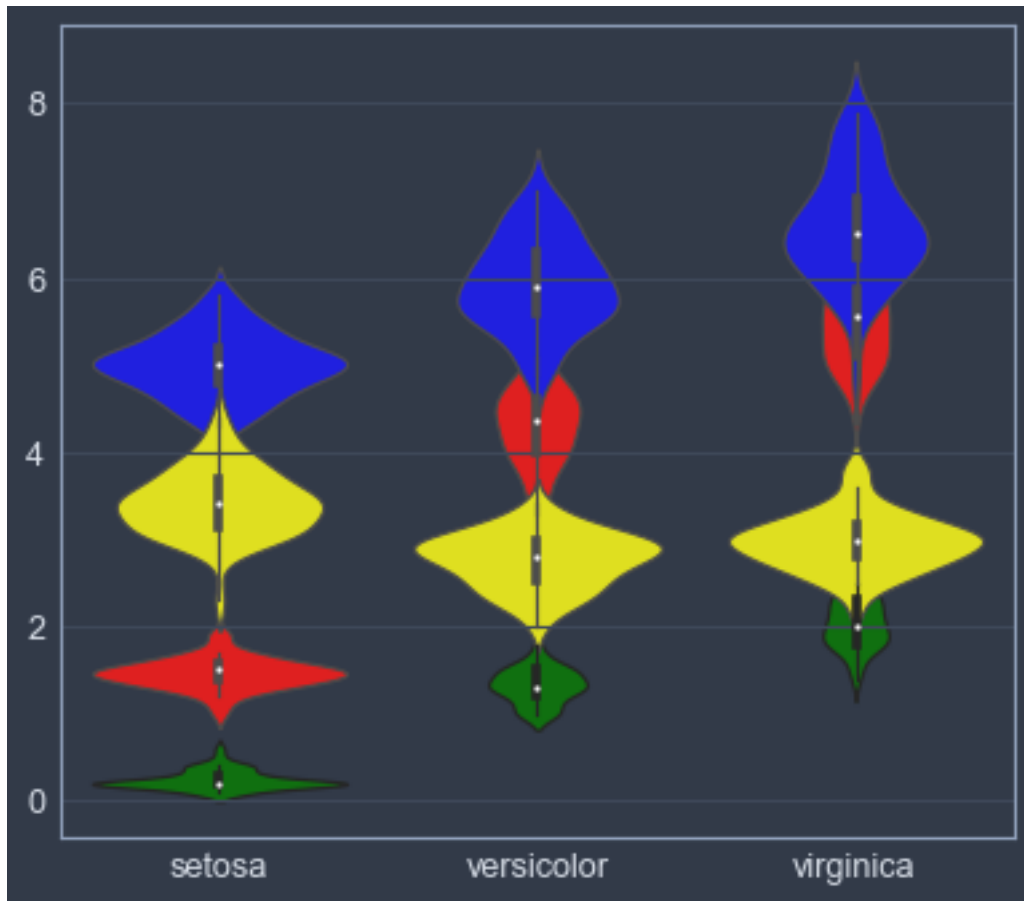
In [22]: sns.swarmplot(x="species", y='petal_length', color='red', data=iris_data)
sns.swarmplot(x="species", y="petal_width", color='green', data=iris_data)
sns.swarmplot(x="species", y="sepal_length", color='blue', data=iris_data)
sns.swarmplot(x="species", y="sepal_width", color='yellow', data=iris_data)
plt.xlabel(''); plt.ylabel('')

Out[22]: Text(0,0.5,'')
```

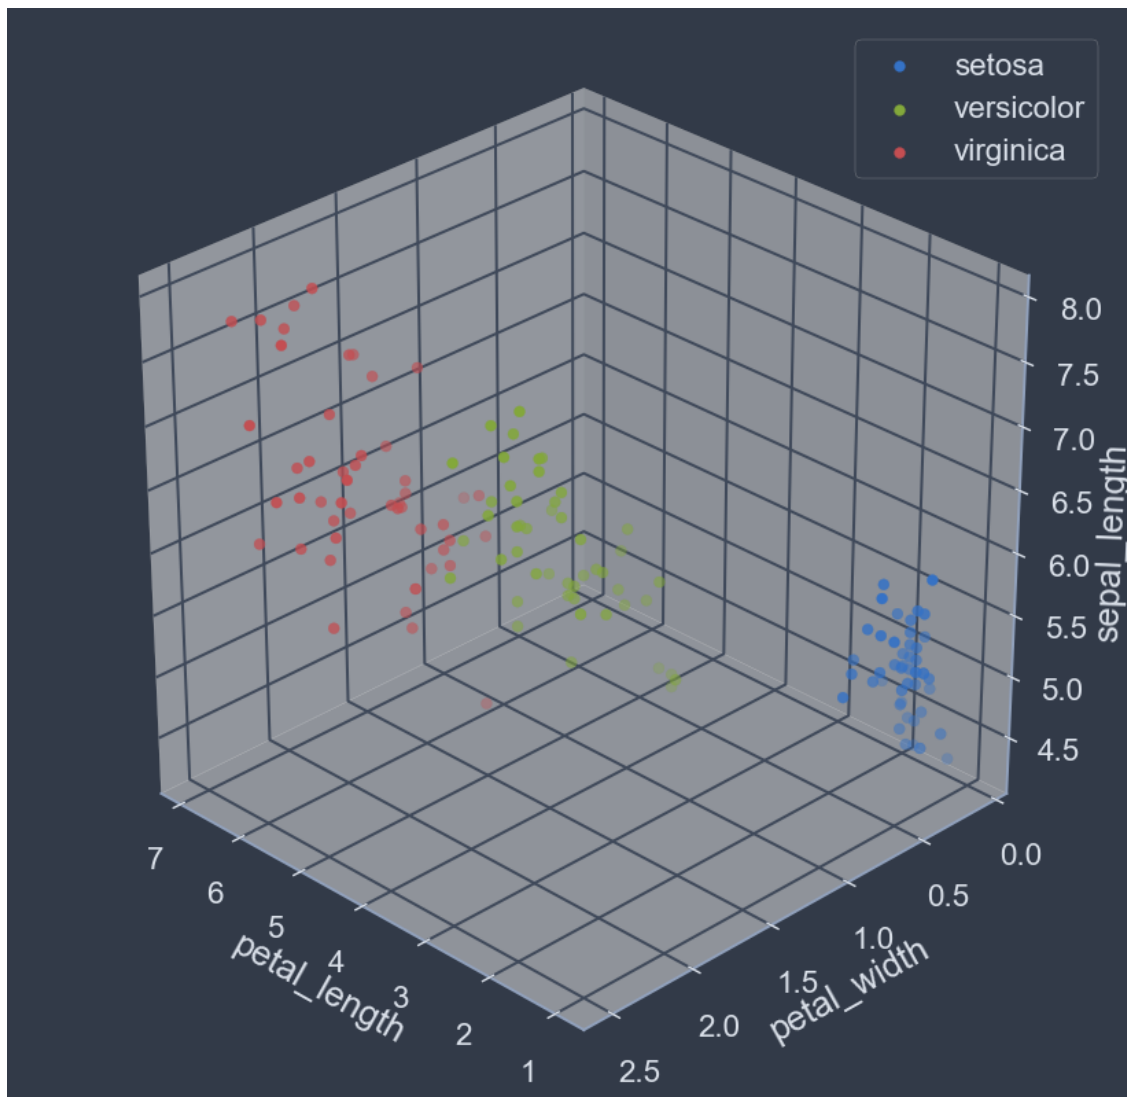


```
In [13]: sns.violinplot(x="species", y='petal_length', color='red', data=iris_data)
sns.violinplot(x="species", y="petal_width", color='green', data=iris_data)
sns.violinplot(x="species", y="sepal_length", color='blue', data=iris_data)
sns.violinplot(x="species", y="sepal_width", color='yellow', data=iris_data)
plt.xlabel(''); plt.ylabel('')
```

```
Out[13]: Text(0,0.5,'')
```



```
In [14]: fig= plt.figure(figsize=(8,8), dpi=128); ax = fig.add_subplot(111, projection='3d')
x_name = 'petal_length'; y_name = 'petal_width'; z_name='sepal_length'
for species in pd.unique(iris_data.species):
    x = iris_data[iris_data.species==species][x_name]
    y = iris_data[iris_data.species==species][y_name]
    z = iris_data[iris_data.species==species][z_name]
    ax.scatter(x,y,z,label=species)
ax.view_init(30, 135)
ax.set_xlabel(x_name); ax.set_ylabel(y_name); ax.set_zlabel(z_name)
plt.legend()
plt.show()
```



Exploration Notes

- The class 'setosa' is linearly well separated.
- The classes 'versicolor' and 'virginica' are likely not linearly separable.
- Classes seem to have good boundaries overall.
- Data forms some well defined clusters.
- Classes are mostly easily distinguishable by looking at cluster membership.

1.2.3 Classifier Selection

Judging by our exploratory analysis, not all classes can be completely separated via a hyperplane. However, the classes do seem to have well defined boundaries and are generally easily distinguishable. Therefore, some form of non-linear classification approach will likely yield better results: commonly by space partitioning via curved hyper-surface class boundaries.

Both SVMs (non-linear, via a feature space transform) and k-NN can be suitable for the task. However, due to the shape of the data, using an SVM will require an extensive hyper-parameter search as well as kernel comparison against various metrics, with no clear foreseeable benefits. Meanwhile, the data distribution looks well suited for a kNN to be quickly and easily trained without much tuning or computation.

By way of the above rationale, we'll choose to utilize the kNN algorithm.

1.2.4 Classifier Training

```
In [119]: # separate data into features and labels
```

```
    X = iris_data.drop(columns='species')
    y = iris_data.species
    print('features: ', X.shape); print('labels: ', y.shape)
```

```
features: (150, 4)
```

```
labels: (150,)
```

```
In [120]: # training validation split
```

```
    X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.1)
    print('train split: ', X_train.shape, y_train.shape)
    print('test split: ', X_test.shape, y_test.shape)
```

```
train split: (135, 4) (135,)
```

```
test split: (15, 4) (15,)
```

```
In [121]: # randomization
```

```
    X_train, y_train = shuffle(X_train, y_train)
```

```
In [122]: # normalization
```

```
    sc = StandardScaler()
    sc.fit(X_train)
    X_train_norm = sc.transform(X_train)
    X_test_norm = sc.transform(X_test)
```

```
In [123]: # choose k-nearest-neighbors param
```

```
    k = int(round(np.sqrt(X_train.shape[0])))
    k
```

```
Out[123]: 12
```

```
In [124]: # train classifier
```

```
    knn = KNeighborsClassifier(n_neighbors=k, metric='euclidean')
    knn.fit(X_train_norm, y_train)
```

```
Out[124]: KNeighborsClassifier(algorithm='auto', leaf_size=30, metric='euclidean',
                                metric_params=None, n_jobs=1, n_neighbors=12, p=2,
                                weights='uniform')
```

1.2.5 Classifier Validation

TODO: ROC, AUC, metrics, k values, visualize boundaries, SVM comparison, plot all 6 scatters, make 3d code more compact

```
In [125]: # score classifier
          print('Train Accuracy: {:.2f}'.format(knn.score(X_train_norm, y_train)))
          print('Test Accuracy: {:.2f}'.format(knn.score(X_test_norm, y_test)))
```

Train Accuracy: 0.96

Test Accuracy: 1.00

1.3 MNIST Dataset

[link to MNIST dataset source website](#)

1.3.1 Overview

The dataset consists of 60,000 labeled samples. Each sample is a grayscale 28x28 pixel PNG image. Each image is labeled as a single digit from the range [0,9].

Note: this is not the original dataset in neither shape, form, format. We've done some pre-processing to it (outside scope of this notebook) for convenience.

1.3.2 Dataset Exploration

```
In [3]: # MNIST image directory and image file extension
        IMG_DIR = './mnist-images/'; IMG_FILE_EXTENSION = '.png'
        # load file paths sorted by numerical value of file name, excluding the extension suffix
        fpaths = sorted(os.listdir(IMG_DIR), key=lambda x: int(x[:-len(IMG_FILE_EXTENSION)]))
        # load images
        X = np.array([cv2.imread(IMG_DIR+fp, cv2.IMREAD_GRAYSCALE) for fp in fpaths], dtype=np.uint8)
        # load labels
        y = np.loadtxt('mnist_labels.txt', dtype=np.uint8)
```

```
In [4]: # count and shape
        print('X: ', X.shape, '\ny: ', y.shape)
```


X: (60000, 28, 28)

y: (60000,)

```
In [5]: # unique classes
        np.unique(y)
```


Out[5]: array([0, 1, 2, 3, 4, 5, 6, 7, 8, 9], dtype=uint8)

```
In [201]: # showcase a few samples
          for ax in plt.subplots(4,8)[1].ravel():
              rnd = np.random.randint(0,y.shape[0])
              ax.set_axis_off()
              ax.imshow(X[rnd], cmap='gray')
```





```
In [202]: # plot class sample counts  
plt.hist(y, bins=range(11))
```


```
Out[202]: (array([ 5923.,  6742.,  5958.,  6131.,  5842.,  5421.,  5918.,  6265.,  
                  5851.,  5949.]),  
array([ 0,  1,  2,  3,  4,  5,  6,  7,  8,  9, 10]),  
<a list of 10 Patch objects>)
```



```
In [203]: # plot logarithmic grayscale color histogram  
hist = plt.hist(X.flatten(), bins=range(256), log=True)
```



```
In [6]: # compute and show average pixel values for each digit class
        avgs = np.hstack([np.mean(X[y==i], axis=0) for i in range(10)])
        plt.matshow(avgs,cmap='gray'); plt.axis('off');
```



```
In [7]: # compute and show median pixel values for each digit class
        medians = np.hstack([np.median(X[y==i], axis=0) for i in range(10)])
        plt.matshow(medians,cmap='gray'); plt.axis('off');
```



```
In [8]: # compute and show standard diviation of pixel values for each digit class
stds = np.hstack([np.std(X[y==i], axis=0) for i in range(10)])
plt.matshow(stds,cmap='gray'); plt.axis('off');
```



1.3.3 Preprocessing

```
In [256]: # training validation split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3)
print('train split: ', X_train.shape, y_train.shape)
print('test split: ', X_test.shape, y_test.shape)
```

```
train split: (42000, 28, 28) (42000,)
test split: (18000, 28, 28) (18000,)
```

```
In [257]: # randomization
X_train, y_train = shuffle(X_train, y_train)
```

```
In [258]: # threshold and binarize the images
X_train = np.where(X_train>127,1,0).astype(np.bool)
X_test = np.where(X_test>127,1,0).astype(np.bool)
plt.matshow(np.hstack(X_train[:10]*255),cmap='gray'); plt.axis('off');
```



```
In [259]: # flatten images to match classifier, preprocessors, etc expected input shapes
X_train = X_train.reshape(len(X_train), -1)
X_test = X_test.reshape(len(X_test), -1)
print('X_train:', X_train.shape, ' X_test:', X_test.shape)
```

```
X_train: (42000, 784) X_test: (18000, 784)
```

```
In [260]: # standarize and perform PCA dimensionality reduction
pca = PCA(0.65, whiten=True)
X_train = pca.fit_transform(X_train)
X_test = pca.transform(X_test)
print('X_train:', X_train.shape, ' X_test:', X_test.shape)
```

```
X_train: (42000, 34) X_test: (18000, 34)
```

```
In [261]: # plot inverse PCA result
plt.matshow(np.hstack(pca.inverse_transform(X_train[:10])).reshape(-1,28,28)), cmap='g
```



```
In [262]: # tsne = TSNE(n_components=2, n_iter=250)
# tsne.fit(X_train)
```

1.3.4 Classifier Selection

MNIST is a larger dataset with many more dimensions and training samples (~5m even in binary form). Although it's still not large enough to make kNN classification computationally unfeasible, it's probably a better direction to utilize something like a linear or polynomial SVM since it's likely to be a much more efficient classifier.

1.3.5 Classifier Training

```
In [271]: # setup SVM hyperparam search
hyperparams = {'kernel':['linear', 'poly', 'rbf'], 'C':expon(scale=10), 'gamma':expon(
rand_search = RandomizedSearchCV(SVC( ), hyperparams, n_jobs=-1)
```

```
In [272]: %%capture
rand_search.fit(X_train,y_train)
```

1.3.6 Classifier Validation

```
In [273]: # rand_search.best_params_
# rand_search.best_estimator_
# rand_search.score(X_test, y_test)
y_pred = rand_search.predict(X_test)
print(classification_report(y_test,y_pred))

precision    recall  f1-score   support
```

0	0.99	0.99	0.99	1771
1	0.99	0.99	0.99	1970
2	0.97	0.97	0.97	1812
3	0.96	0.97	0.97	1807
4	0.96	0.98	0.97	1749
5	0.97	0.96	0.96	1589
6	0.98	0.99	0.98	1804
7	0.97	0.97	0.97	1874
8	0.97	0.96	0.97	1786
9	0.97	0.96	0.97	1838
avg / total	0.97	0.97	0.97	18000